

Object Recognition

Shayeree Sarkar, Baishali Mullick*

*Department of Electrical and Computer Engineering, Carnegie Mellon University

18752 - Estimation, Detection and Learning
Course Project, Spring 2020

1 Problem Statement

Object recognition is breaking into a wide range of industries, with use cases ranging from personal security to productivity in the workplace including image retrieval, security, surveillance, automated vehicle systems and machine inspection. Significant challenges stay on the field of object recognition including, how to train models that generalize well to real-world settings that have not been seen in training and how to better exploit small-scale training data. While deep learning has shown great success in various tasks with a large amount of labeled data, current techniques generally break down if few labeled examples are available.

2 Train and Test Data Collection and Data Visualisation

We have used the Cifar-10 dataset for the purpose of object recognition and collected the preprocessed images for training from <https://www.cs.toronto.edu/~kriz/cifar.html>
We collect random test data images from google and preprocess it using normalization and data augmentation with the Keras data framework.¹

3 Methods

3.1 Logistic Regression

Binomial and Multinomial Logistic Regression are popular classifiers used especially on Discriminative Models

We perform the following steps :

1. We load the data set and split it into train, validation and test dataset using masking
2. We preprocess the images for the three split datasets by normalizing the images, and then reshaping them from the original form of (32,32,3) to (3073,1) for each image after adding a bias unit

*Email: shayeres@andrew.cmu.edu, bmullick@andrew.cmu.edu

¹The data visualisation using PCA is depicted in the appendix section 4.1

25 3. For binary logistic regression we form a subset of two classes out of the ten classes, namely,
26 cats and dogs by randomly sampling images for these classes from the train dataset and their cor-
27 responding labels

28 4. We train the weight vector using gradient descent algorithm for **200 epochs** and **learning rate**
29 of **0.0001** and achieve a **training accuracy** of **56 percent** using cost function in 3.2, and update
30 parameters using equation 3.4

31 5. We use the sigmoid function for the prediction of the labels such that ones **greater than 0.5** are
32 classified as either cats or dogs and we achieve a **test accuracy** of **51 percent**

33 6. For Multinomial Logistic Regression we train the weight vector using **stochastic gradient**
34 **descent** algorithm this time, for a subset of 20000 images containing equal samples for each class
35 and train it for **150 epochs** only to achieve achieve a **training accuracy of 80 percent** and **test**
36 **accuracy of 68 percent** with the **learning rate** of **0.0001** using cost function in 3.2, and update
37 parameters using equation 3.4 ²

38 3.2 Convolutional Neural Network

39 3.2.1 VGG-3 Baseline Model

40 We have explored CNN design with a **VGG-3 architecture**. The results of the model on the test
41 dataset showed an improvement in classification accuracy with each increase in the depth of the
42 model. Overfitting starts at around 15-to-20 epochs. We have used **Dropouts, L2 Regularization**
43 **factor**, and **Data Augmentation** of the training images to achieve this accuracy with **learning rate**
44 of **0.001** after **125 epochs** with a Batch size of 64 which has helped in achieving an accuracy of **89**
45 **percent** on the training dataset and **88 percent** on test dataset ³

46 3.2.2 VGG-3 Baseline Model with SVM

47 Conventionally, the Softmax function is the classifier used at the last layer of the CNN network.
48 Usage of linear support vector machine (SVM) in the last layer of the CNN instead, often has
49 been proven to give high test accuracy. We have used **Dropouts, regularization factor, and data**
50 **augmentation** of the training images to achieve this accuracy with **learning rate** of **0.001** after
51 **125 epochs** with a Batch size of 64 and the loss as **Hinge Loss** helping in achieving an accuracy of
52 **85.1 percent** on the **training dataset** and **83.4 percent** on **test dataset**. ⁴

53 3.3 Support Vector Machines using Linear Kernel

54 The objective of the support vector machine algorithm is to find a hyperplane in an N-dimensional
55 space(N —the number of features) that distinctly classifies the data points by margin maximization.
56 The plot of accuracy against C for the linear kernel is shown above. The performance of the SVM
57 on both the test set and training set improve substantially until **C = 0.1**, after which the performance
58 diminishes. The best **test accuracy of 28.6 percent** was obtained for **C = 0.1**. We conclude SVM
59 alone is a **poor classifier** for image data. ⁵

²All graphs and output image data is plotted in Appendix Section 4.2.1

³All graphs and output image data is plotted in Appendix Section 4.2.2

⁴All graphs and output image data is plotted in Appendix Section 4.2.3

⁵All graphs and output image data is plotted in Appendix Section 4.2.4

60 4 Appendix

61 4.1 Train and Test Data Collection and Data Visualisation

The data visualisation using PCA for CIFAR-10 dataset is depicted below:

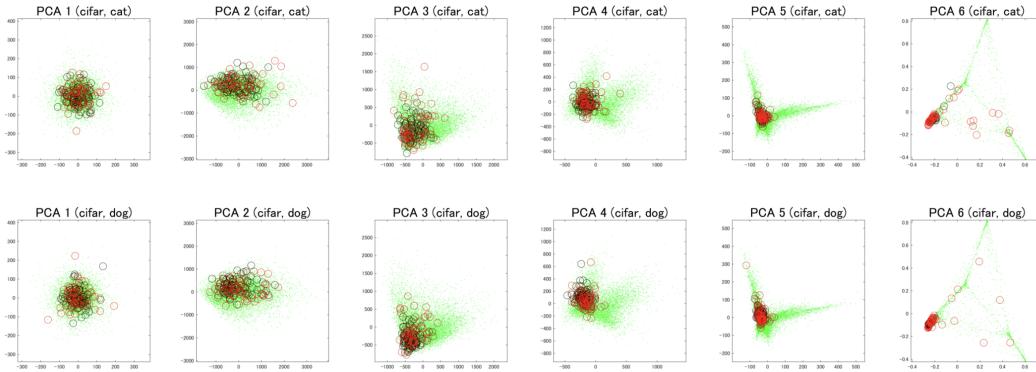


Figure 1: Visualization results for CIFAR-10 data (“cat” and “dog” classes) by usual PCA

62

63 4.2 Methods

64 4.2.1 Logistic Regression

65 Binomial and Multinomial Logistic Regression are popular classifiers used especially on Discriminative Models They use the following activation function :

$$p(y \mid x = 1, \Theta) = \frac{\exp(\theta_y x)}{\sum_{j=1}^K \exp(\theta_j x)} = \text{softmax}((\Theta x)_y) \quad (4.1)$$

67 We perform the following steps :

- 68 1. We load the data set and split it into train, validation and test dataset using masking
- 69 2. We preprocess the images for the three split datasets by normalizing the images, and then
70 reshaping them from the original form of **(32,32,3)** to **(3073,1)** for each image after adding a bias
71 unit
- 72 3. For binary logistic regression we form a subset of two classes out of the ten classes, namely,
73 cats and dogs by randomly sampling images for these classes from the train dataset and their cor-
74 responding labels
- 75 4. We train the weight vector using gradient descent algorithm for **200 epochs** and **learning rate**
of **0.0001** and achieve a **training accuracy of 56 percent** using cost function in 3.2, and update
parameters using equation 3.4

$$J(\theta) = -\log p(\mathbf{|X}, \theta) = \sum_{i=1}^N -y^{(i)} (\theta^T x^{(i)}) + \log \left(1 + e^{\theta^T x^{(i)}} \right) \quad (4.2)$$

75 5. We use the sigmoid function for the prediction of the labels such that ones **greater than 0.5** are
76 classified as either cats or dogs and we achieve a **test accuracy of 51 percent**

6. For Multinomial Logistic Regression we train the weight vector using **stochastic gradient descent** algorithm this time, for a subset of 20000 images containing equal samples for each class and train it for **150 epochs** only to achieve achieve a **training accuracy of 80 percent** and **test accuracy of 68 percent** with the **learning rate of 0.0001** using cost function in 3.2, and update parameters using equation 3.4

$$\frac{\partial J(\theta)}{\partial \theta_j} = - \sum_{i=1}^N \begin{pmatrix} y^{(i)} \\ 1 \end{pmatrix} \left[y^{(i)} - \frac{e^{\theta^T(i)}}{1 + e^{\theta^T(i)}} \right] \quad (4.3)$$

$$\theta_j \leftarrow \theta_j - \eta \frac{\partial J(\theta)}{\partial \theta_j}, (\eta = learningrate) \quad (4.4)$$

77 Results for Logistic Regression

The graphs we obtain are shown below:

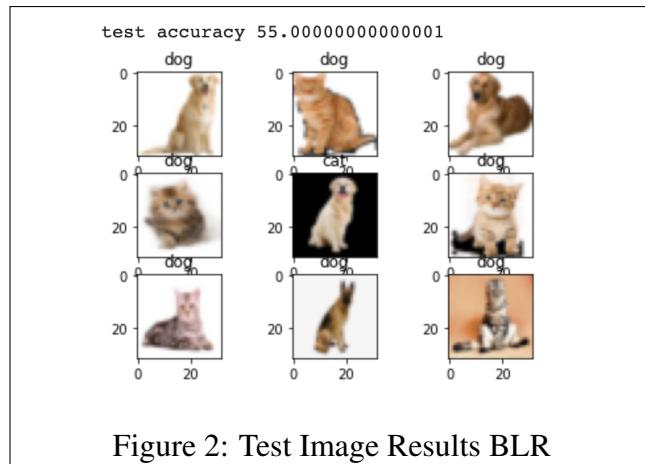


Figure 2: Test Image Results BLR

78

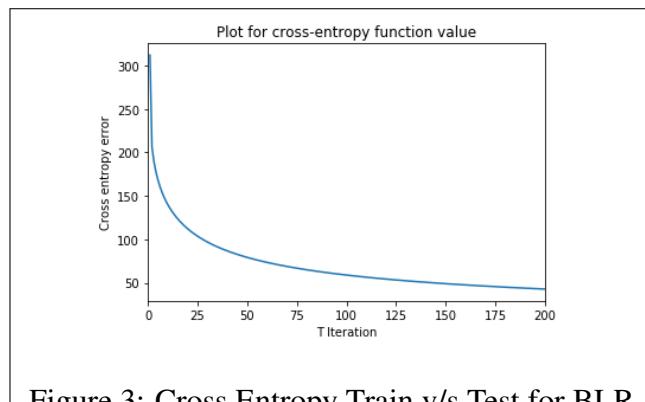


Figure 3: Cross Entropy Train v/s Test for BLR

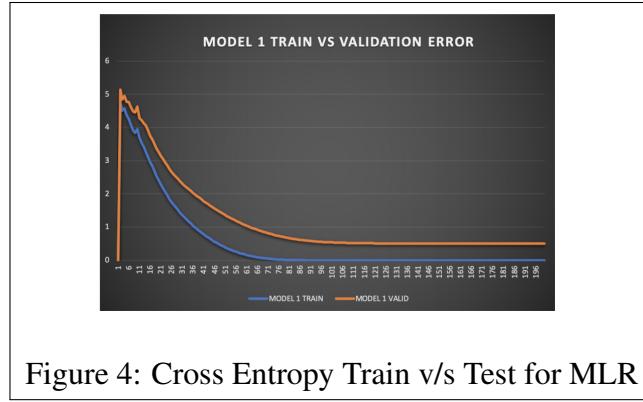


Figure 4: Cross Entropy Train v/s Test for MLR

79 4.2.2 Convolutional Neural Network

- 80 VGG-3 Baseline Model We have explored CNN design with a **VGG-3 architecture**. The results of
 81 the model on the test dataset showed an improvement in classification accuracy with each increase
 82 in the depth of the model. Overfitting starts at around 15-to-20 epochs.
 83 We have used **Dropouts**,**L2 Regularization factor**, and **Data Augmentation** of the training im-
 84 ages to achieve this accuracy with **learning rate** of **0.001** after **125 epochs** with a Batch size of 64
 85 which has helped in achieving an accuracy of **89 percent** on the training dataset and **88 percent**
 on test dataset **Results for CNN:**

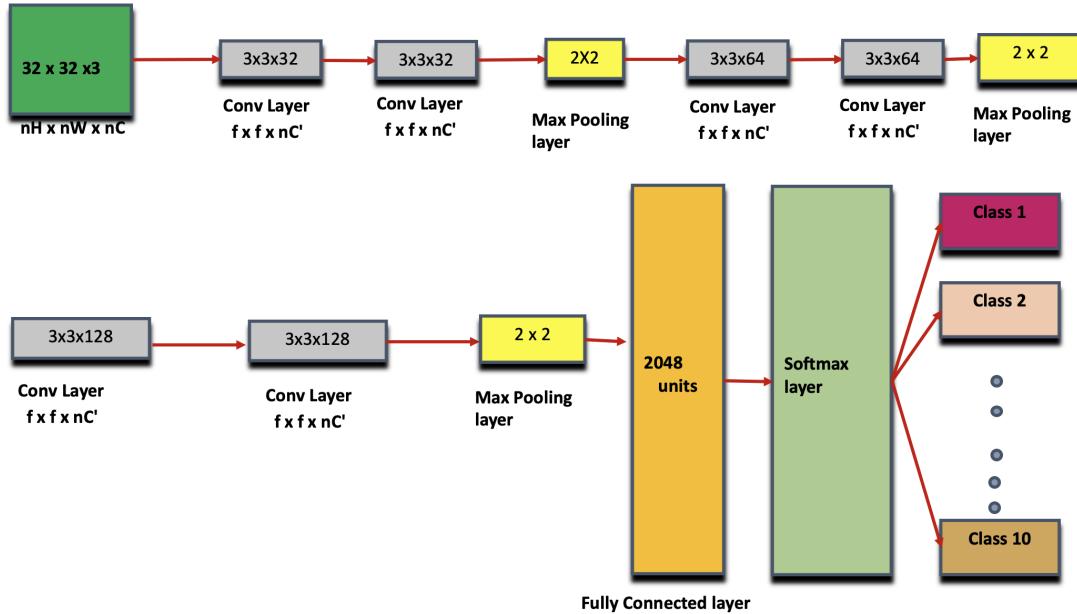
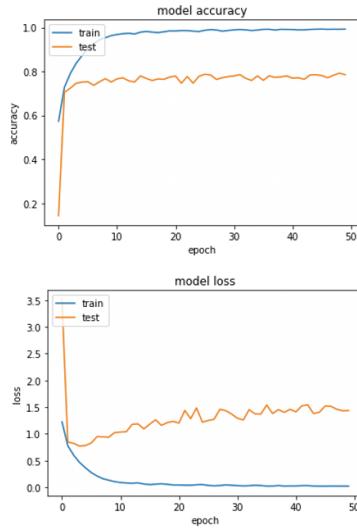


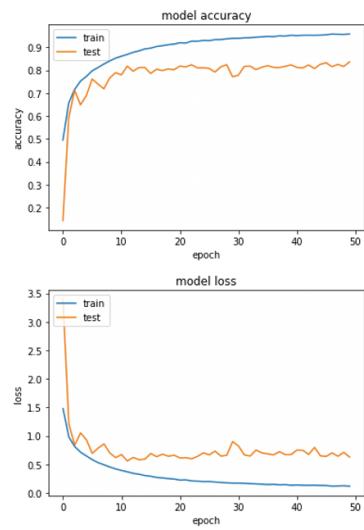
Figure 5: VGG-3 Baseline Model Architecture

Training and test loss and accuracy for CNN without regularization and 256 dense units after flattening in 50 epochs



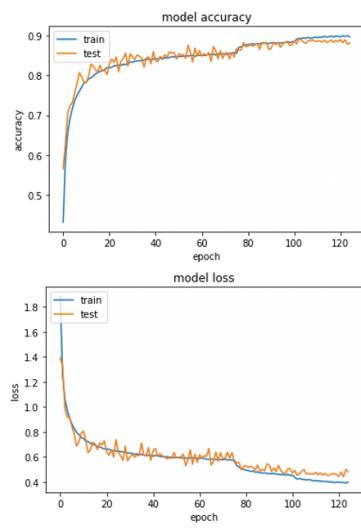
Train Accuracy : 99.8%
Test Accuracy : 78.5%

Training and test loss and accuracy for CNN without regularization and 512 dense units after flattening in 50 epochs



Train Accuracy : 95.84%
Test Accuracy : 83.7%

Training and test loss and accuracy for CNN with regularization and no dense units after flattening in 125 epochs



Train Accuracy : 89%
Test Accuracy : 88.1%

Figure 6: VGG-3 Baseline Model Architecture Model Comparisons

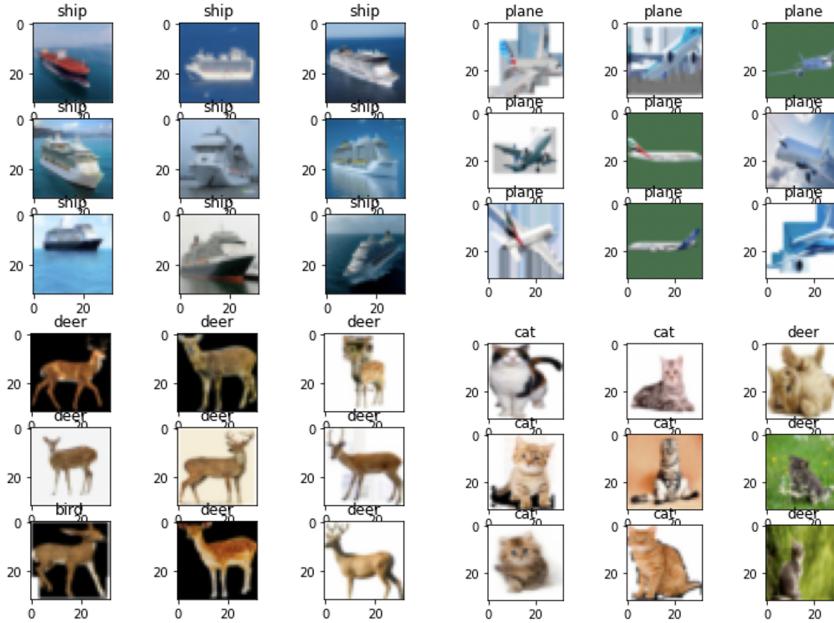


Figure 7: CNN Output with VGG-3 Baseline Model on Test Images

87 4.2.3 CNN VGG-3 Baseline Model with SVM

88 Conventionally, the Softmax function is the classifier used at the last layer of the CNN network.
89 Usage of linear support vector machine (SVM) in the last layer of the CNN instead, often has been
90 proven to give high test accuracy.

91 We have used **Dropouts, regularization factor, and data augmentation** of the training images to
92 achieve this accuracy with **learning rate of 0.001** after **125 epochs** with a Batch size of 64 and the
93 loss as **Hinge Loss** helping in achieving an accuracy of **85.1 percent** on the **training dataset** and
94 **83.4 percent** on **test dataset**. We can see the Architecture in Fig 8, the loss and accuracy plots in
95 Fig 9 and results on test images in Fig 10.

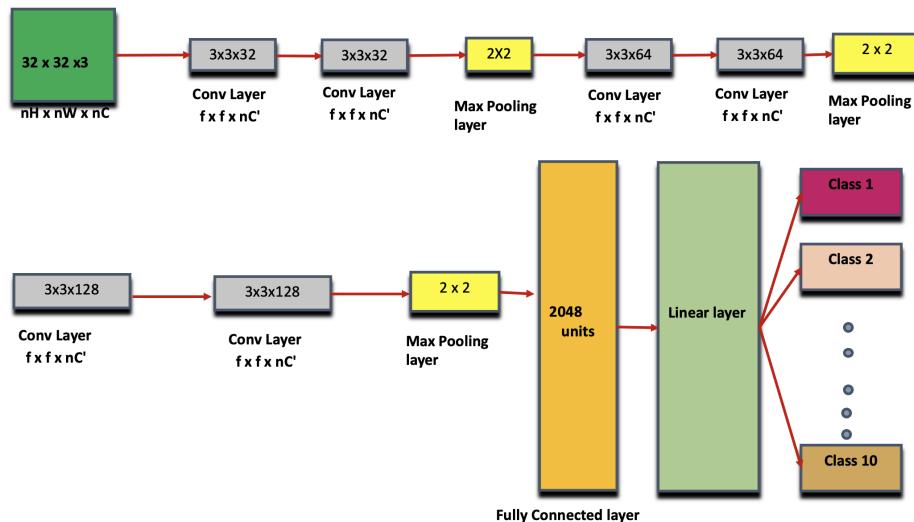


Figure 8: CNN with VGG 3 baseline Model with SVM

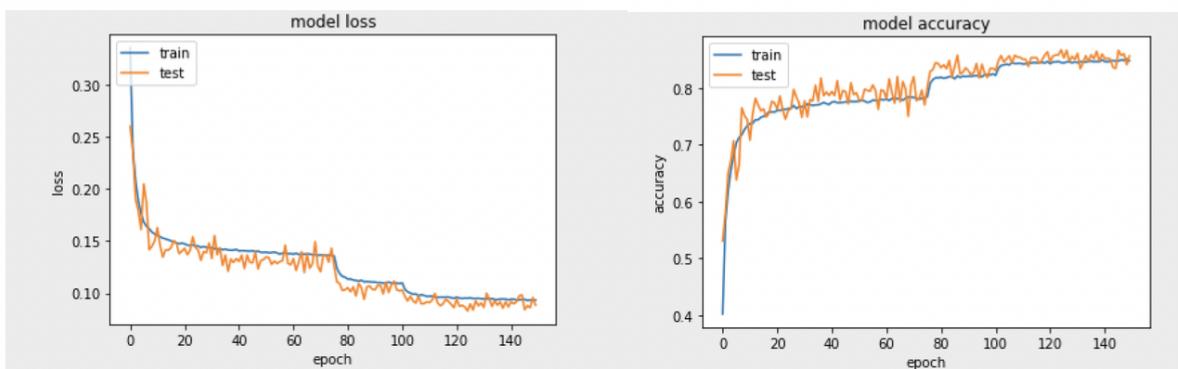


Figure 9: Model accuracy and Loss for Train v/s Test for CNN with SVM

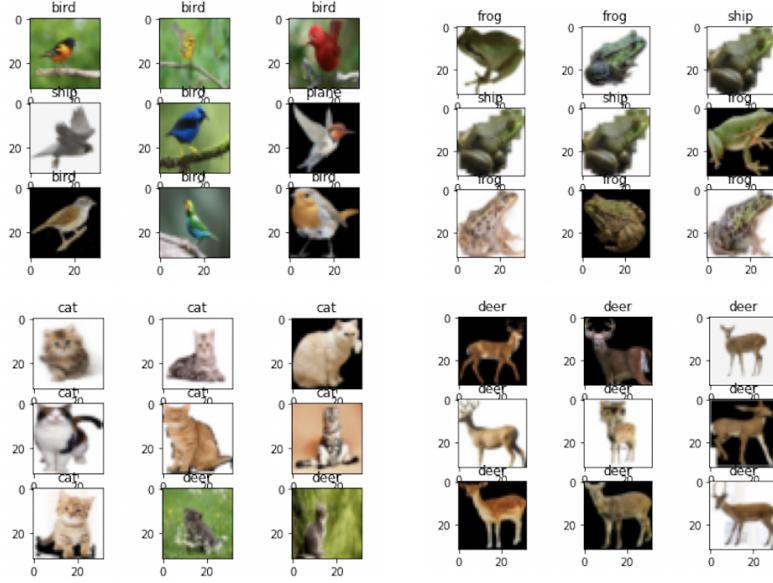


Figure 10: Ouput on Test Images of CNN with VGG 3 baseline Model with SVM

96 4.2.4 Support Vector Machines using Linear Kernel

97 The objective of the support vector machine algorithm is to find a hyperplane in an N-dimensional
 98 space(N —the number of features) that distinctly classifies the data points by margin maximization.
 99 The plot of accuracy against C for the linear kernel is shown above. The performance of the SVM
 100 on both the test set and training set improve substantially until **C = 0.1**, after which the performance
 101 diminishes. The best **test accuracy of 28.6 percent** was obtained for **C = 0.1**. We conclude SVM
 102 alone is a **poor classifier** for image data. Image Classification on the CIFAR-10 Dataset using
 103 Support Vector Machines is done using Linear Kernel. The results are shown in Fig 11 and Fig 12.

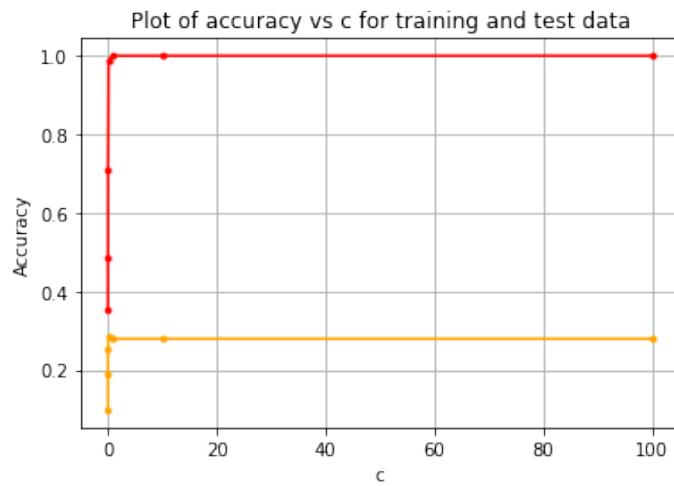


Figure 11: Train v/s Test Accuracy with SVM

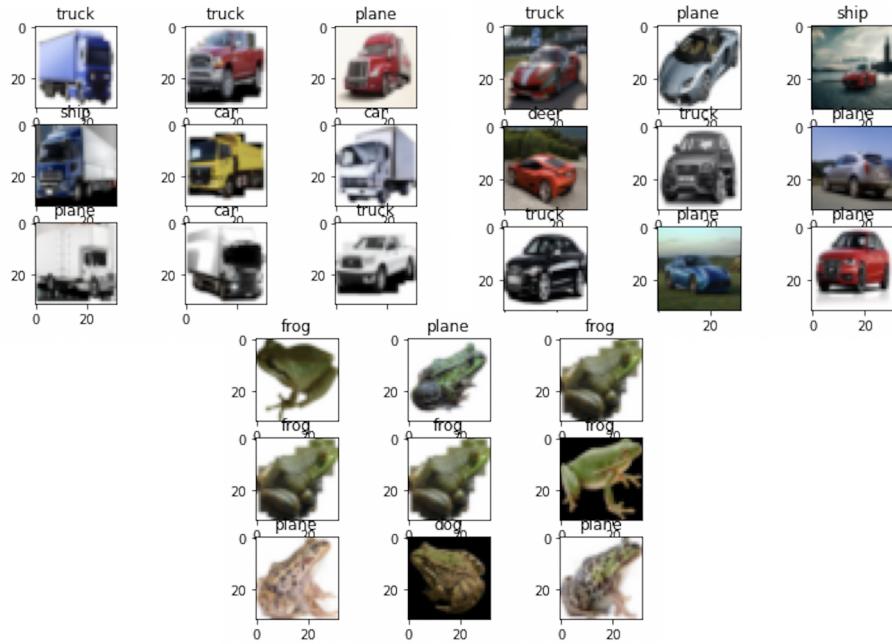


Figure 12: Output for Test Images with SVM

104 **5 Acknowledgments**

105 Shayeree Sarkar and Baishali Mullick acknowledges the help and support of Professor Rohit Negi
106 and the Teaching Assistant of the course 18752 - Zinan Lin